

COMP1411 (Spring 2022) Introduction to Computer Systems

Individual Assignment 2

Duration: 00:00, 19-Mar-2022 ~ 23:59, 20-Mar-2022

<i>Name</i>	
<i>Student number</i>	

Question 1. [3 marks]

In this question, we use the Y86-64 instruction set (please refer to Lecture 4-6).

1(a) [1 mark]

Write the machine code encoding of the assembly instruction:

`"mrmovq 0x230(%rax), %rcx"`.

Please write the bytes of the machine code in hex-decimal form, i.e., using two hex-decimal digits to represent one byte. You are allowed to leave spaces between adjacent bytes for better readability. The machine has a little-endian byte ordering.

Show your steps. Only giving the final result will NOT get a full mark of this question.

Answer:

(a)

1) `mrmovq`'s op code: 50

2) `rA = %rcx`, `rB = %rax`, so the second byte will be 10

3) displacement = 0x230, written into 8-byte little-endian encoding is: 30 02 00 00 00 00 00 00

4) the final machine code: 50 10 30 02 00 00 00 00 00 00

1(b) [2 marks]

Consider the execution of the instruction “`mrmovq 0x230(%rax), %rcx`”. Assume that for now, the data in register `%rax` is 0x100, just before executing this instruction, the value of PC is 0x300. Please use “**vm**” to represent the actual data reading from the main memory.

Describe the steps done in the following stages: Fetch, Decode, Execute, Memory, Write Back, PC update, by filling in the blanks in the table below.

Note that you are required to fill in the generic form of each step in the second column; and in the third column, fill in the steps for the instruction “`mrmovq 0x230(%rax), %rcx`” with the above given values. If you think there should not be a step in some stage, just leave the blanks unfilled.

The symbol “ \leftarrow ” means reading something from the right side and assign the value to the left side. X:Y means assign the highest 4 bits of a byte to X, and assign the lowest 4 bits of the byte to Y.

Answer:

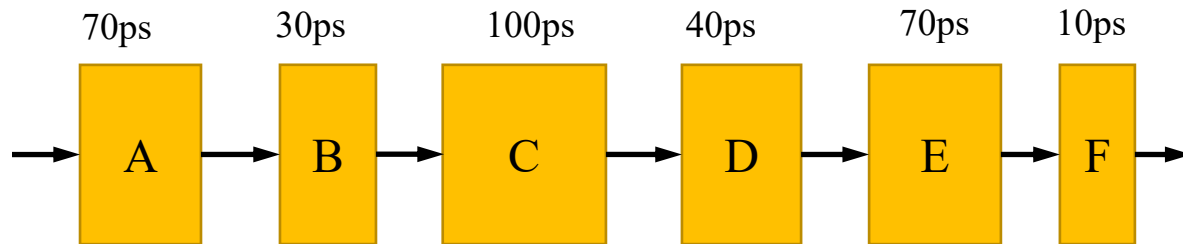
Stages	<code>mrmovq D(rB), rA</code>	<code>mrmovq 0x230(%rax), %rcx</code>
Fetch	<code>icode: ifun \leftarrow _M₁[PC]_</code> <code>rA:rB \leftarrow _M₁[PC+1]_</code> <code>valC \leftarrow _M₈[PC+2]_</code> <code>valP \leftarrow _PC+10_</code>	<code>icode: ifun \leftarrow _M₁[0x300]=5:0_</code> <code>rA:rB \leftarrow _M₁[0x301]=1:0_</code> <code>valC \leftarrow _M₈[0x302]=0x230_</code> <code>valP \leftarrow _0x300+10=0x30A_</code>
Decode	<code>valB \leftarrow _R[rB]_</code>	<code>valB \leftarrow _R[%rax]=0x100_</code>
Execute	<code>valE \leftarrow _valB + valC_</code>	<code>valE \leftarrow _0x100+0x230=0x330_</code>
Memory	<code>valM \leftarrow M₈[_valE]</code>	<code>valM \leftarrow M₈[0x330] = vm</code>
Write back	<code>R[rA] \leftarrow __valM_</code>	<code>R[%rcx] \leftarrow _vm_</code>
PC update	<code>PC \leftarrow _valP_</code>	<code>PC \leftarrow _0x30A_</code>

GRADING SCHEME:

- (1) Total marks: 3 marks, 1(a) 1 mark, and 1(b) 2 marks.
- (2) For question 1(a), correctly writing out steps 1) and 2) and 3) get 0.2 marks each; correctly writing out the machine code in the required hex-decimal format gets 0.4 marks.
- (3) For question 1(b), the answer for each of column 2 and column 3 has 1 mark. For each column, correctly filling each line gets 0.1 marks. Correctly filling all lines gets another 0.1 marks.

Question 2. [3 marks]

Suppose a combinational logic is implemented by 6 serially connected components named from A to F. The whole computation logic can be viewed as an instruction. The number on each component is the time delay spent on this component, in time unit ps, where $1\text{ps} = 10^{-12}$ second. Operating each register will take 20ps.



Throughput is defined as how many instructions can be executed on average in one second for a pipeline, and the unit of throughput is IPS, instructions per second.

Latency refers to the time duration starting from the very first component and ending with the last register operation finished, the time unit for latency is ps.

For throughput, please write the result in the form $X.XX * 10^Y$ IPS, where X.XX means one digit before the dot and two fractional digits after the dot, and Y is the exponent.

2(a) Making the computation logic a 3-stage pipeline design that has the maximal throughput. Note that a register shall be inserted after each stage to separate their combinational logics. [1.5 marks]

- Please answer how to partition the stages.
- Please compute the throughput and latency for your pipeline design, with steps.

Answer:

3 registers inserted: between B and C, between C and D, and after F

Throughput: $1 / ((120 + 20) * 10^{-12}) = 7.14 * 10^9$ IPS

Latency: $(120 + 20) * 3 = 420$ ps

2(b) Making the computation logic a 4-stage pipeline design that has the maximal throughput. Note that a register shall be inserted after each stage to separate their combinational logics. [1.5 marks]

- Please answer how to partition the stages.
- Please compute the throughput and latency for your pipeline design, with steps.

4 registers inserted: between B and C, between C and D, between D and E, and after F

Throughput: $1 / ((100 + 20) * 10^{-12}) = 8.33 * 10^9$ IPS

Latency: $(100 + 20) * 4 = 480$ ps

GRADING SCHEME:

- (1) Total marks: 3 marks, 2(a) 1.5 marks, and 2(b) 1.5 marks.
- (2) For either 2(a) or 2(b), correctly specifying how the stages are partitioned gets 0.5 marks; correctly giving the throughput gets 0.5 marks; correctly giving the latency gets 0.5 marks.

Question 3. [4 marks]

The following byte sequence is the machine code of a program function compiled with the Y86-64 instruction set (refer to Lecture 6). The memory address of the first byte is 0x300. Note that the byte sequence is written in hex-decimal form, i.e., each number/letter is one hex-decimal number representing 4 binary bits, and two numbers/letters represent one byte.

**630030F3020000000000000030F11E000000000000007023030000000
00000601061316211761F03000000000000090**

Please write out the assembly instructions (in Y86-64 instruction set) corresponding to the machine codes given by the above bytes sequence, and explain what this program function is computing.

Answer:

0x300: xorq %rax, %rax

0x302: irmovq \$2, %rbx

0x30C: irmovq \$30, %rcx

0x316: jmp .L2

L1:

0x31F: addq %rcx, %rax

0x321: subq %rbx, %rcx

L2:

0x323: andq %rcx, %rcx

0x325: jg .L1

0x32E: ret

The program computes: $30 + 28 + 26 + \dots + 2$

GRADING SCHEME:

- (1) Total marks: 4 marks
- (2) Correctly giving the assembly instructions gets 3 marks, and correctly explaining the behavior of the program gets 1 mark.
- (3) Specifically, there are 9 instructions, correctly writing out each instruction gets 0.3 marks, correctly writing out all instructions gets another 0.3 marks. Note that if the jumping target for any jump instruction is incorrect, the corresponding instruction is deemed incorrect. If you forget \$ symbol in specifying immediate numbers, which is a

syntax error, you will not get 0.2 marks for the instruction. Students are not required to write out the addresses of the instructions.