

CS3402 Database Systems

Assignment 2 (due: 21/Nov/2022 @11:59pm!)

Question A.

- 1) What are the advantages and disadvantages of hash functions relative to B-tree indices? How might the type of index available influence the choice of a query-processing strategy? **[15 marks]**

Answer:

Hash functions have a few advantages over B-tree:

- requires no indices, hence no extra space cost (3 marks)
- can support selection-type operations very fast (3 marks)

Hash functions have some disadvantages relative to B-tree:

- collisions must be handled by hash function look-ups (3 marks)
- being non-clustering, range search will be very slow (3 marks)

Thus, depending on the type of query, if selection-type operations are involved, we may choose hashing if it is available, or if range queries are to be processed, we may go for B-tree if it's available. (3 marks)

- 2) When is it preferable to use a dense index rather than a sparse index? **[8 marks]**

Answer:

It is preferable to use a dense index rather than a sparse index when

- the file is not ordered on the indexed field (such as when the index is a secondary index) (4 marks)
- or, when the index file is small enough to fit in the main memory (4 marks)

- 3) What is the difference between primary index and a secondary index? **[7 marks]**

Answer:

The difference between primary index and a secondary index lies in the facts that:

- The primary index is on the field which specifies the sequential order of the data file (4 marks)
- There can only be one primary index while many secondary indices (3 marks)

- 4) If a hash structure is used on a search key for which range queries are likely, what property should the hash functions have? **[10 marks]**

Answer:

The hash function should preserve the order of the data.

That is, if h is the hash function, and $x < y$, then it implies that $h(x) < h(y)$.

Question B.

Consider the relations:

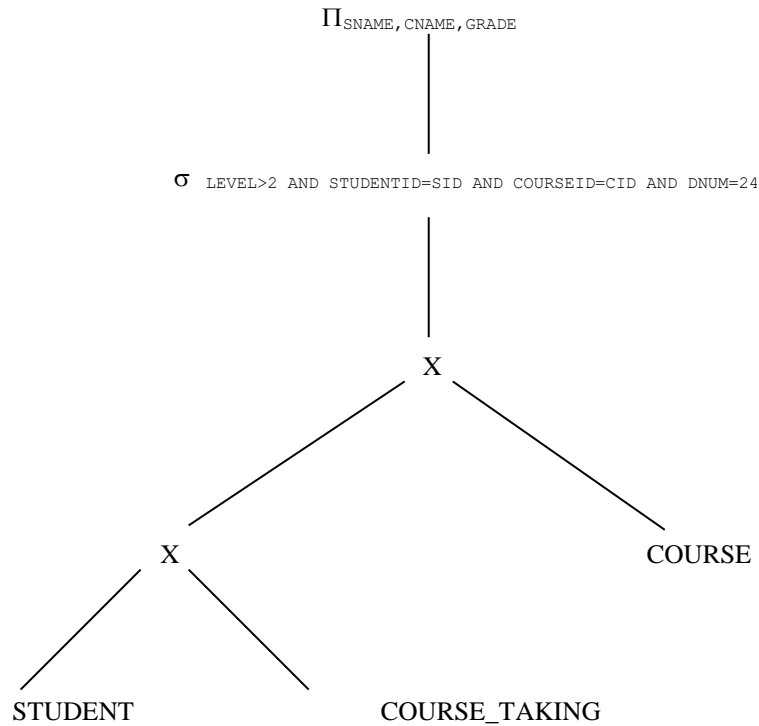
STUDENT(SNAME, SID, BDATE, ADDRESS, DNUM)
COURSE(CNAME, CID, LEVEL, LECTURER)
COURSE_TAKING(STUDENTID, COURSEID, GRADE)

as well as the following SQL query:

```
SELECT SNAME, CNAME, GRADE
FROM STUDENT, COURSE_TAKING, COURSE
WHERE LEVEL > 2 AND STUDENTID = SID
      AND COURSEID = CID AND DNUM = 24;
```

- a) Draw a canonical query tree for the above SQL query. **[9 marks]**
- b) Apply the optimization rules to the above query tree and come up with the most optimized query tree using those rules. State the necessary assumptions for your decision. **[21 marks]**


Answer:



Marking scheme: partially correct gets =<3 points.

- b) Apply the optimization rules to the above query tree and come up with the most optimized query tree using those rules. State the necessary assumptions for your decision. **(21 marks)**

Assumption: *There are less courses with level > 2 , compared to the number of students in department 24.*



Marking scheme: One optimization step missing will get 2 points deducted and some typos or missing symbols will also cause 2-4 points deducted.

Question D.

[30 marks]

- 1) Do you agree that a nonrecoverable schedule results in a loss of transaction atomicity? Explain your answer.
- 2) When the system recovers from a crash, in what order must transactions be undone and redone?
- 3) Suppose that the system crashes during the time it is recovering from a prior crash. When the system again recovers, what action must be taken?

Answer:

1) Yes: If a failure occurs, it may be necessary to *abort a committed transaction*, because in a nonrecoverable schedule a transaction T is allowed to commit *before* all the transactions from which T has read some data have committed. **(8 marks)**

Answer:

2) First, undo all transactions that began prior to the crash but have not committed (ie, there is a **start** record but no **commit** record in the log). These transactions are undone *in reverse temporal order* (newest to oldest); Next, redo all transactions that committed prior to the crash (ie, there is a **commit** record in the log). These transactions are redone *in temporal order* (from oldest to newest). **(14 marks)**

Answer:

3) The answer is *exactly the same* as for that of Question D 2) above, because undo and redo are idempotent (so it is acceptable to undo or to redo a transaction several times). The database will be consistent once the recovery activity has been completed. **(8 marks)**